

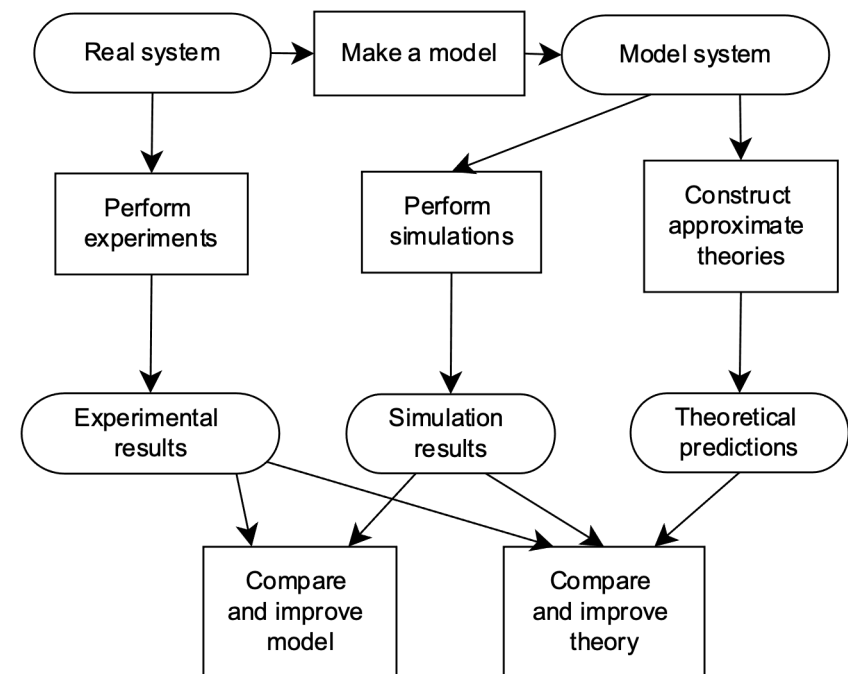
Geospatial Data Science
Content Block II: *Techniques*
Lab 6
Generate & Analyze Spatial RVs

Austin J. Brockmeier, Ph.D.

Wednesday, March 15th, 2023

Probability and Spatial Statistics for Geospatial Systems and Data

1. Statistical analysis
2. Generate data from random variables and processes through **computer simulation** (Lab 6&7)
3. Analyze spatial patterns in data



Outline

- Generating random numbers
- Counting: binning/histograms
- Basic statistics
- Bayes rule: joint and conditional
- Spatial correlation



Coin data




Binary data: Flip a fair coin, create a **Bernoulli** random variable X

$$X(\omega) = \begin{cases} +1, \omega = \text{Heads} \\ 0, \omega = \text{Tails} \end{cases}$$

$$P_X(1) = \Pr(X = 1) = p$$
$$P_X(0) = \Pr(X = 0) = 1 - p$$

$$X \in \{0, 1\}$$




Count data:

Flip a fair coin N times, count of heads is **binomial** random variable S with probability mass function

$$P_S(n) = \binom{n}{N} p^N (1 - p)^{n-N}$$

$$S \in \{0, 1, \dots, N\}$$



Analysis Questions

What is the proportion of heads if we observe 10, 100, or 1000 coin flips?

What is mean value?

What are the confidence intervals for the proportion of heads?

Random number generation

Pseudo random number generators (deterministic if seed is known):


https://en.wikipedia.org/wiki/Pseudorandom_number_generator

Generating normal RVs from uniform:

https://en.wikipedia.org/wiki/Box%E2%80%93Muller_transform

```
import numpy as np
import scipy.stats as st

rng = np.random.default_rng(seed=0)
```



Coin data



[scipy.stats.bernoulli](https://docs.scipy.org/doc/scipy/reference/stats/bernoulli.html)

```
X = st.bernoulli.rvs(p, size=n_trials)
print(X)
print("{} flips with {} heads ({}%)".
      format(n_trials,np.sum(X==1),np.mean(X==1)))
```

```
1000 flips with 519 heads (0.519%)
```

```
[1 0 0 0 1 0 0 0 1 0 1 1 0 1 1 1 0 0 1 1 1 1 1 1 0 0 1 1 0 1 1 0 1 1 1
1 0 0 0 0 1 0 1 1 0 0 1 1 1 1 0 1 0 1 1 1 0 0 0 0 1 1 0 1 1 0 0 0 1 1 0 0 1 0 1
0 1 0 1 1 1 0 1 0 1 0 0 0 0 0 1 1 0 0 1 0 0 1 0 0 0 1 1 1 0 1 1 0 1 0 0 0 0 1 0
0 1 1 0 1 0 1 1 0 0 1 1 0 0 1 1 0 1 1 0 0 1 1 0 1 1 1 1 0 1 1 0 0 0 0 0 1 0
0 1 0 0 0 1 0 0 1 1 1 0 1 1 1 0 1 1 1 0 1 0 1 0 1 1 1 1 1 1 0 0 0 0 0 1 0
0 1 0 0 0 1 0 0 0 1 1 0 0 1 1 0 0 0 1 1 1 1 1 0 1 0 1 1 1 1 1 1 0 1 0 1 1
0 0 0 0 1 1 1 1 1 0 1 1 1 1 1 0 1 0 1 0 0 1 1 0 0 0 0 0 1 1 0 0 0 0 0 0 1
0 1 1 1 0 1 0 0 0 0 0 1 0 0 1 0 0 0 1 0 0 0 1 0 1 1 0 0 0 0 1 1 1 0 1 0 1
1 0 0 1 0 1 0 1 1 0 0 0 1 0 0 0 1 1 1 1 0 1 0 0 1 0 0 0 1 1 1 1 0 0 0 1 0
1 1 0 0 1 1 1 1 0 0 0 0 0 0 0 1 0 1 0 0 1 1 0 0 0 1 0 1 1 1 1 0 0 0 0 1
1 0 0 1 0 0 0 0 0 1 0 1 1 1 0 1 1 0 1 1 0 0 0 0 1 0 0 1 1 0 1 1 1 1 0 0 0
1 1 0 1 0 0 1 0 0 1 1 0 1 1 0 0 0 0 0 0 1 0 0 1 1 1 0 1 1 1 0 1 1 1 1 0
0 1 1 1 0 0 1 0 1 0 0 1 1 1 0 0 0 1 0 1 1 1 0 1 0 1 0 1 0 1 0 1 0 1 0 1
1 1 1 1 1 0 1 1 0 0 1 0 0 1 0 1 0 1 0 1 1 0 0 1 1 0 0 1 1 0 0 1 0 0 0 1
1 1 0 1 1 1 0 1 0 1 0 1 1 0 0 0 0 0 1 0 1 1 1 1 1 0 0 1 1 1 1 1 1 0 0 1
0 0 1 0 1 0 0 0 0 1 0 0 1 0 0 1 1 0 0 1 1 1 0 0 0 1 0 1 1 0 0 1 0 1 0 0 0
0 0 0 0 1 0 0 1 0 0 1 1 1 0 1 0 1 1 1 0 0 0 1 1 1 0 1 1 1 0 1 0 1 0 0 1
1 0 1 0 1 1 0 1 0 0 1 1 0 1 0 1 1 1 0 1 0 1 0 0 0 1 1 1 1 0 1 1 0 0 0 0 0
0 0 1 0 0 0 0 0 0 1 0 1 0 1 1 0 1 1 1 1 1 1 1 1 0 0 0 1 0 1 0 0 0 0 1 1
0 0 0 0 0 0 0 0 1 0 0 0 1 1 1 0 0 1 1 1 0 0 0 1 0 0 1 0 0 0 0 0 0 1 0 0
1 1 1 0 0 0 1 0 0 0 1 1 0 1 0 0 1 1 1 0 1 0 1 0 0 1 1 1 1 1 0 1 1 1 1 0
0 0 1 1 1 0 1 0 0 1 1 1 1 0 1 0 0 0 1 1 1 0 1 0 0 1 0 0 0 0 0 1 1 1 0 0 0
1 1 0 0 1 0 1 0 1 0 0 0 0 1 1 0 1 0 0 1 0 0 1 1 0 1 0 0 0 1 0 1 0 1 0 0 1
0 0 0 0 0 0 1 0 1 0 1 1 0 1 1 1 1 0 0 0 1 0 0 0 0 1 1 1 1 1 0 0 0 0 0 1 0
0 0 1 1 0 0 0 1 0 1 1 0 0 1 1 1 1 0 0 1 1 1 0 1 0 1 1 1 0 0 1 1 1 0 1 1
1 0 0 1 1 0 1 1 0 0 0 1 0 0 1 1 1 0 0 0 0 1 1 0 1 0 1 0 1 1 1 1 1 1 1 0
0 1 1 1 0 0 0 1 1 1 0 0 1 0 0 1 1 1 0 0 1 0 1 0 0 0 0 0 0 1 0 1 0 0 0 1
0]
```

Coin data



[numpy.random.Generator.binomial](#)

```
rng = np.random.default_rng(seed=0)
N = 1
n_trials, p = 1000, 0.5
X = rng.binomial(N, p, n_trials)
print("{} flips with {} heads ({}%)".
      format(n_trials, np.sum(X==1), np.mean(X==1)))
```

```
1000 flips with 527 heads (0.527%)
```

Counting in NumPy



numpy.sum numpy.newaxis



```
p = 0.25
n_trials = 4
X = rng.binomial(1, p, n_trials)
L=X[:,np.newaxis]==[0,1]
counts = np.sum(L,axis=0)
```

X

[0 0 0 1]

→

X[:,np.newaxis]

[[0]
[0]
[0]
[1]]

X[:,np.newaxis]==[0,1]

[[0]
[0]
[0]
[1]]

==

[0 1]

→

L

[[True False]
[True False]
[True False]
[False True]]

np.sum(L, axis=0)

→

counts

[3 1]

axis=0

Coin data

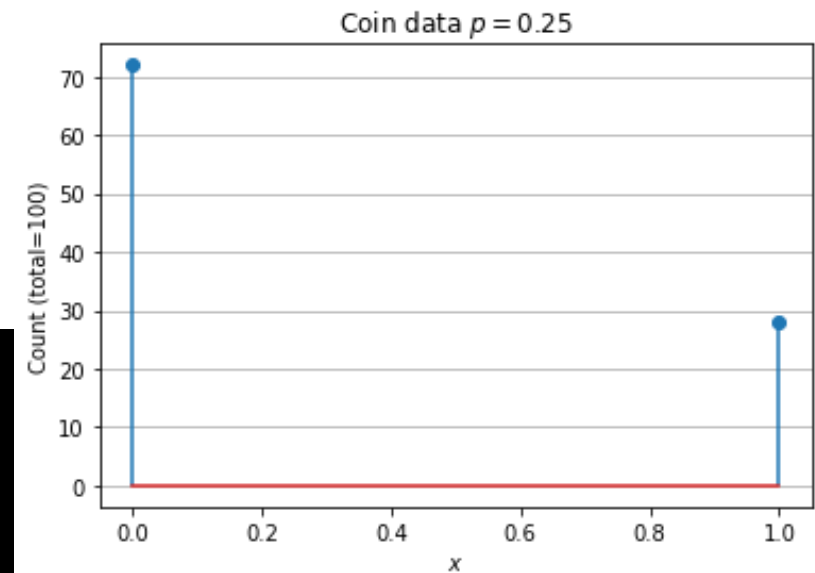


matplotlib.stem



```
p = 0.25
n_trials = 100
X = rng.binomial(1, p, n_trials)
print(X)
plt.figure()
plt.stem([0,1],np.sum(X[:,np.newaxis]==[0,1],axis=0))
plt.ylabel("Count (total={})".format(n_trials))
plt.xlabel(r'$x$')
plt.grid(axis='y')
plt.title('Histogram {}'.format(r"$p=$",p))
```

```
[0 0 0 1 0 0 0 0 1 0 0 0 1 1 0 0 1 0 1 1 0 0 0 0 0 1 0 1 1 0 1 0 0 0 0 0 0 0 0 0 0 1 0
0 0 0 0 0 0 0 0 1 1 0 0 1 0 0 1 0 1 0 0 1 1 0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 0 0 0 0 1
0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1]
```




Coin data



[numpy.histogram](#) or [matplotlib.hist](#)

```
p = 0.25
n_trials = 100
X = rng.binomial(1, p, n_trials)
print(X)
plt.figure()

plt.hist(X)
plt.xlabel(r'$x$')
plt.grid(axis='y')
plt.title('Histogram {}'.format(r'$p=$', p))
```



Not clear what values X takes.... 0.1?
0.9

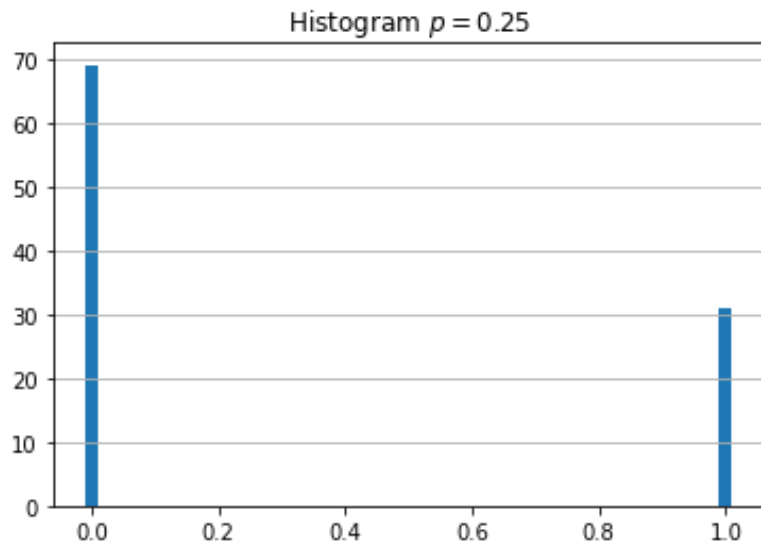
Coin data



matplotlib.hist



```
plt.figure()
plt.hist(X, bins=[-0.01, 0.01, 0.99, 1.01])
plt.xlabel(r'$x$')
plt.grid(axis='y')
plt.title('Histogram {}{} '.format(r"$p=$", p))
```



Notes: bins

All but the last (righthand-most) bin is half-open.

In other words, if *bins* is:

[1, 2, 3, 4]

then the first bin is **[1, 2)** (including 1, but excluding 2) and the second **[2, 3)**. The last bin, however, is **[3, 4]**, which *includes* 4.

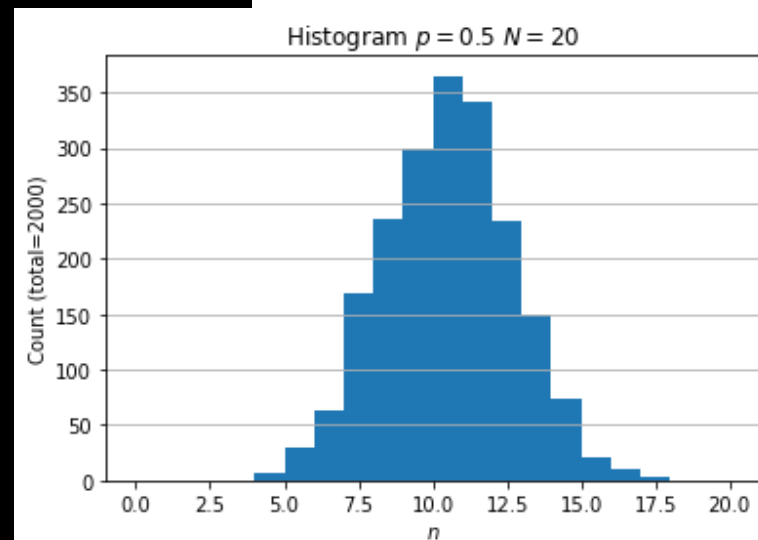
Coin data



[numpy.histogram](#) or `matplotlib.hist`

```
p, N = 0.5, 20
S = rng.binomial(N, p, n_trials)
n_vals = np.linspace(0, N, N+1)

plt.figure()
plt.hist(S, bins=n_vals)
plt.ylabel("Count  
(total={})".format(n_trials))
plt.xlabel(r'$n$')
plt.grid(axis='y')
plt.title('Histogram {}{} {}{}' .
          format(r"$p=$", p, r"$N=$", N) )
```




Coin data



scipy.stats.binom

```
plt.figure()
n_vals = np.arange(0,N)
p_S = st.binom.pmf(n_vals, N, p)
ax = plt.stem(n_vals,p_S)
plt.xlabel(r'$n$')
plt.ylabel(r"$p_S(n)$")
plt.grid(axis='y')
plt.title('{} {} {}'.format(r"$p=$",p,r"$N=$",N))
plt.show()
```



Coin data






numpy.convolve



The probability mass function (pmf) of the sum of two random variables is equal to the convolution of their pmfs

```
def binomial_pmf(p, N):  
    pmf = np.array((1-p, p))  
    pmf_binomial = pmf  
    for i in range(N-1):  
        pmf = np.convolve(pmf, pmf_binomial)  
    return np.linspace(0, N, N+1), pmf
```


```
N, p = 2, 0.25  
plt.figure(figsize=(3, 3))  
n_vals, p_N = binomial_pmf(p, N)  
ax = plt.stem(n_vals, p_N)  
plt.xlabel(r'$n$')  
plt.ylabel(r'$p_S(n)$')  
plt.grid(axis='y')  
plt.title('{} {} '.format(r"$p=$", p, r"$N=$", N))
```



Coin data



```
N, n_trials = 1000, 2000
S = rng.binomial(N, p, n_trials)
plt.plot(S, linestyle='None', marker='x', color='k',
         label="N={}, p={}".format(N, p))
plt.ylabel('Number of heads on trial')
plt.xlabel('trial index')
plt.legend()
```




Coin data



[numpy.quantile](#)

```
quantiles95 = np.array((0.025, 0.975))
conf_interval = np.quantile(S/N, quantiles95)
print("95% confidence interval on prob of heads is  
{0:.4f} {0:.4f}".format(conf_interval))
plt.plot(S/N, linestyle='None', marker='x', color='k')
plt.plot([-1, n_trials], conf_interval[0]*np.ones(2), color='r', linewidth=4)
plt.plot([-1, n_trials], conf_interval[1]*np.ones(2), color='m', linewidth=4)
plt.xlabel('trial index')
plt.ylabel('Number of heads on trial')
```

95% confidence interval on prob
of heads is [0.469 0.529025]




Coin data



The shape of the violin is a density estimate!

[matplotlib.pyplot.violinplot](https://matplotlib.org/3.1.1/api/_as_gen/matplotlib.pyplot.violinplot.html)

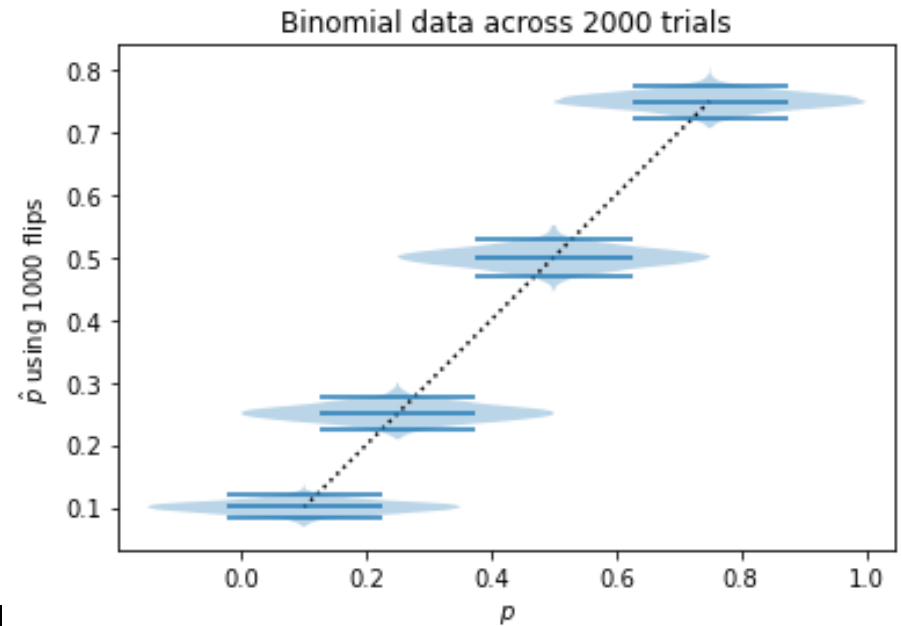
```
p = 0.5
p_hat = np.random.binomial(n_flips, p, N)/N
plt.violinplot(p_hat, showextrema=False, showmeans=True,
               quantiles=quantiles95)
plt.xlabel('$p$')
plt.ylabel('{} using {} flips'.format(r"$\hat{p}$", n_flips) )
plt.title('Binomial data across {} trials'.format(n_trials))
```



Coin data

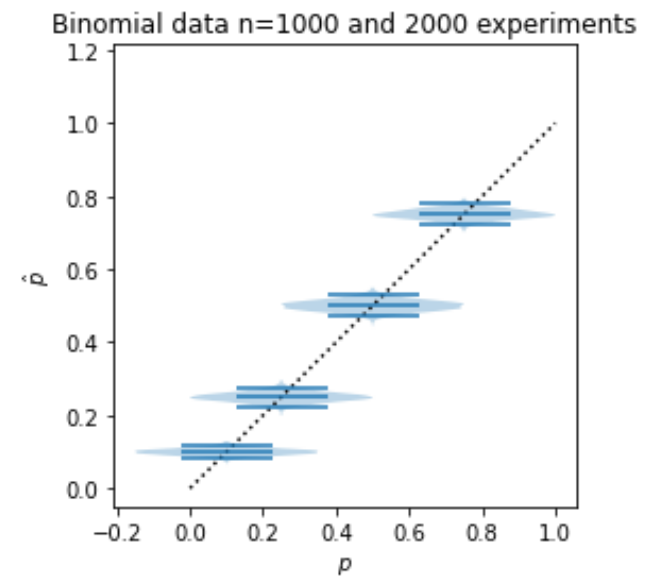

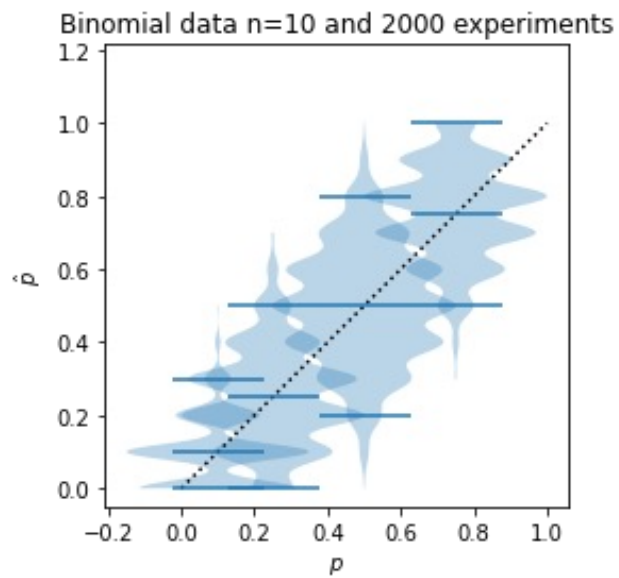


matplotlib.pyplot.violinplot



```
p_list = [0.1, 0.25, 0.5, 0.75]
p_hat_across_p = [rng.binomial(N, p, n_trials)/N for p in p_list]
plt.plot([0,1],[0,1], 'k', linestyle=":")
plt.violinplot(p_hat_across_p, positions=p_list, showextrema=False,
               showmeans=True, quantiles=np.repeat(quantiles95[:,np.newaxis],
               len(p_list), axis=1))
plt.axis('square')
plt.xlabel('$p$')
plt.ylabel('{} using {} flips'.format(r"$\hat{p}$", N) )
plt.title('Binomial data across {} trials'.format(n_trials))
```

Coin data



Estimators: bias, variance, mean squared error

- Expected bias of a parameter estimate:

$$\mathbb{E}[\widehat{\theta(X)}] - \theta^*$$



- Variance of a parameter estimate:

$$\mathbb{E} \left[\left(\widehat{\theta(X)} - \mathbb{E}[\widehat{\theta(X)}] \right)^2 \right]$$

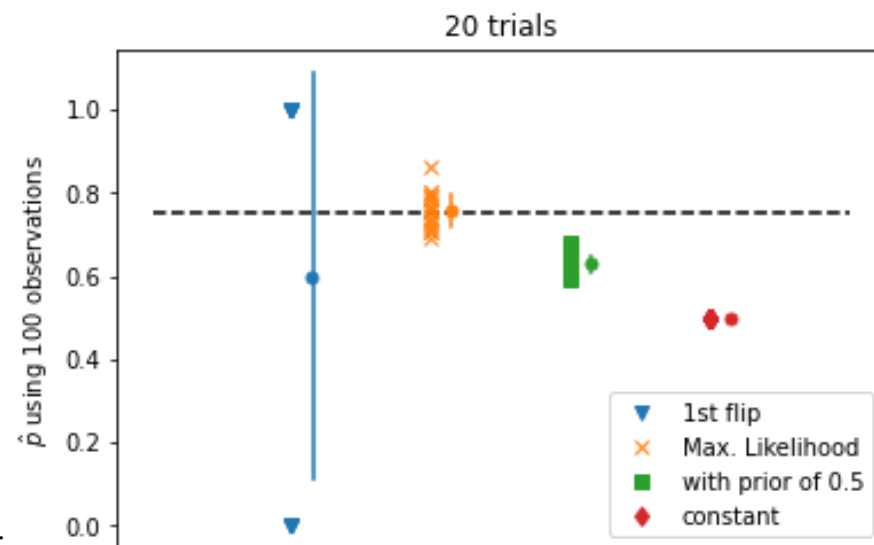
- RMSE of a parameter estimate:

$$\sqrt{\mathbb{E} \left[\left(\widehat{\theta(X)} - \theta^* \right)^2 \right]}$$

Consider an unfair/loading coin. The goal is to estimate the proportion of tails.



- Estimator Silly:** Use first flip to estimate probability of tail.
- Estimator ML:** Observed proportion, this maximizes the likelihood (ML) of the data.
- Estimator Prior:** Weighted average of the observed proportion (weight = N) and 0.5 (weight = 100).
- Estimator Ignorant:** Say the probability of a tail is 0.5 no matter what.



Estimators: bias, variance, mean squared error

Expected bias of a parameter estimate:

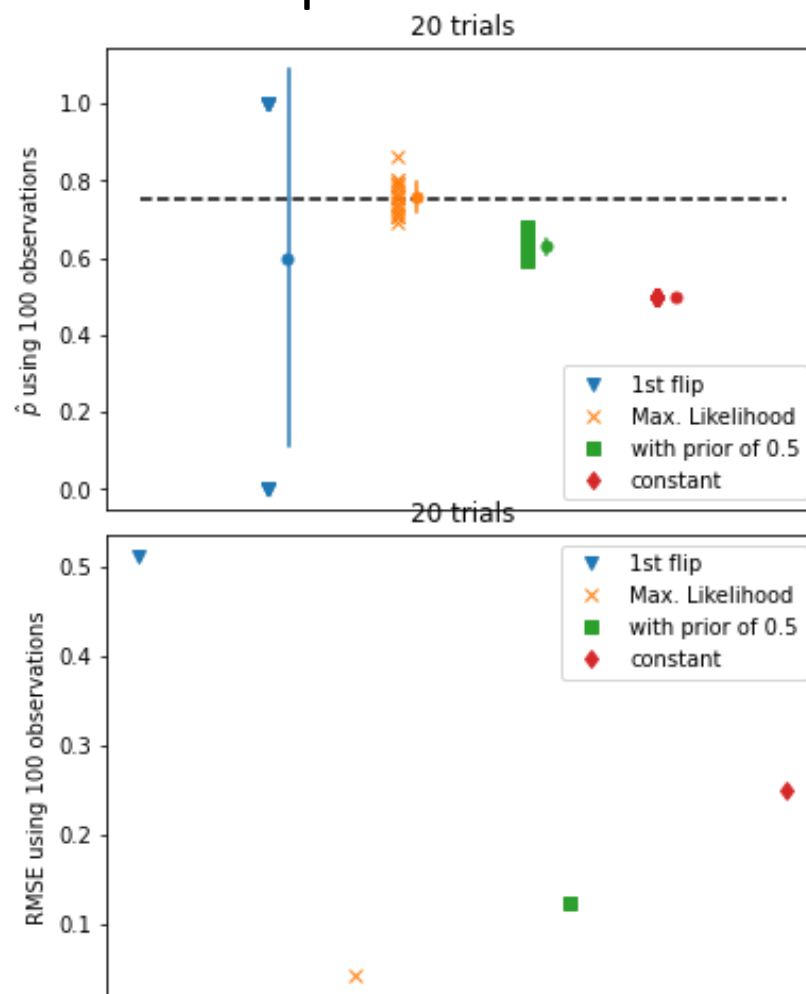
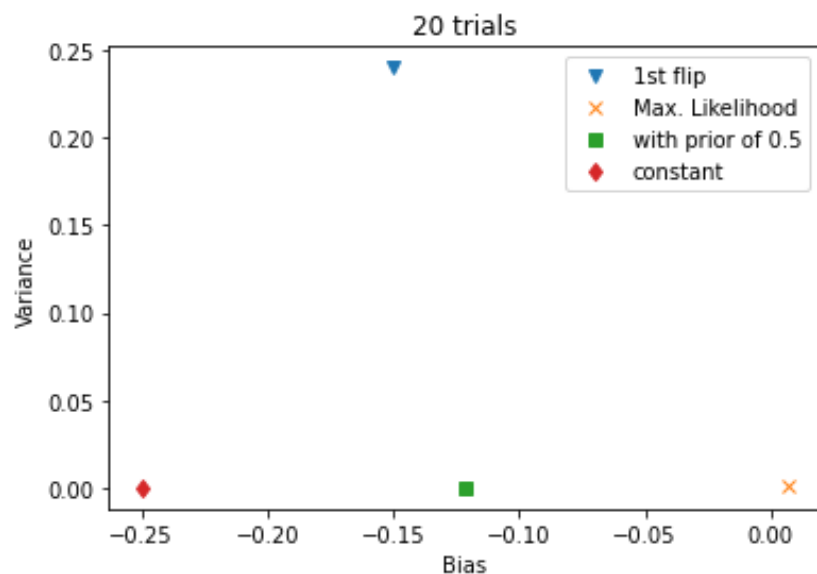
$$\mathbb{E}[\widehat{\theta(X)}] - \theta^*$$

Variance of a parameter estimate:

$$\mathbb{E}[(\widehat{\theta(X)} - \mathbb{E}[\widehat{\theta(X)}])^2]$$

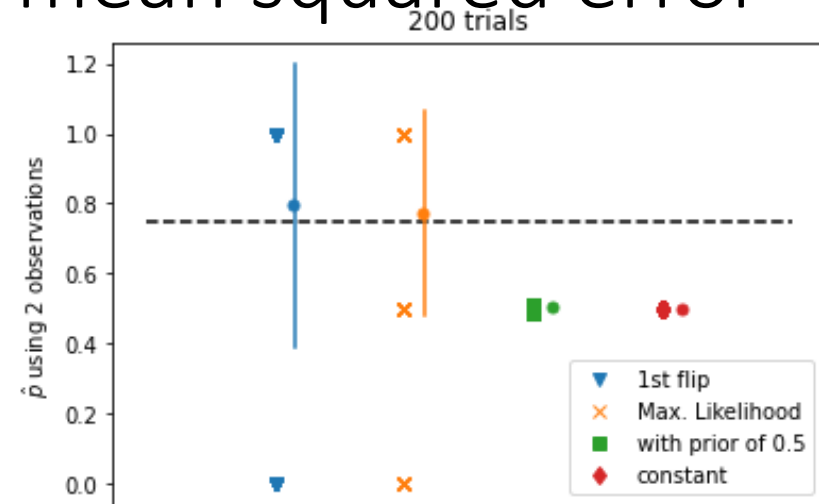

RMSE of a parameter estimate:

$$\sqrt{\mathbb{E}[(\widehat{\theta(X)} - \theta^*)^2]}$$



Estimators: bias, variance, mean squared error

- With only 2 observations ML no longer wins in terms of RMSE.



Coin data



numpy.sum, numpy.ones, numpy.vstack

```
X = bernoulli.rvs(p, size=(n_flips,n_trials))

good_p = np.sum(X==0,axis=0)/n_flips
silly_p = X[0,:]==0
ignorant_p = 0.5*np.ones(n_trials)
prior_weight = 100
prior_value =0.5
prior_p = (n_flips*np.sum(X==0,axis=0)/n_flips
          + prior_weight*prior_value)/(n_flips+prior_weight)

names = ['1st flip',"Max. Likelihood",
        "with prior of 0.5",'constant']
estimates = np.vstack([silly_p,good_p,prior_p,ignorant_p])
```

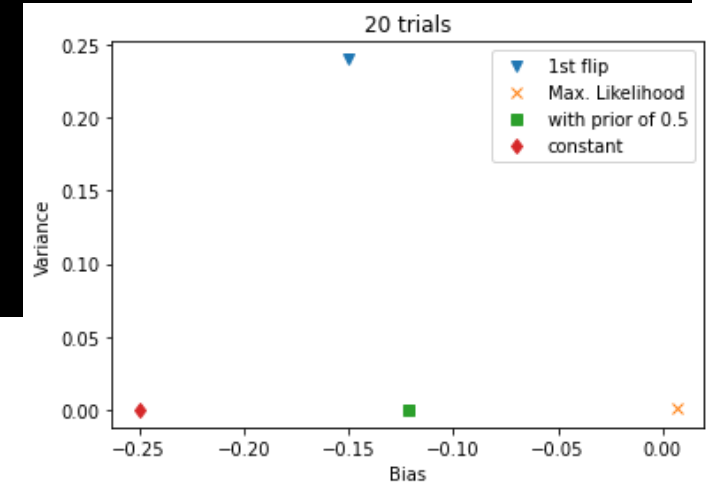
Coin data



numpy.mean, numpy.var, numpy.std

```
markers = ['v', 'x', 's', 'd', '+']
plt.figure()
for i, estimate_name in enumerate(names):
    bias = np.mean( estimates[i])-(1-p)
    h = plt.plot(bias, np.var(estimates[i]), marker=markers[i],
                 linestyle='None', label=estimate_name)

plt.xlabel('Bias')
plt.ylabel('Variance')
plt.title('{} trials'.format(n_trials))
plt.legend(loc='best')
```



Coin data




Maximum likelihood estimation



scipy.stats.binom.pmf

```
def bernoulli_likelihood(ps, X):  
    like = []  
    for p in ps:  
        like.append(st.binom.pmf(np.sum(X), len(X), p))  
    return like
```

```
ps = np.linspace(0,1,1000)  
  
N, p = 100, 0.25  
X = st.bernoulli.rvs(p, size=N)  
plt.plot(ps, bernoulli_likelihood(ps, X), 1  
         label='{} tails out of {} from p={}'.format(np.sum(X), N, p))  
plt.xlabel(r'$\hat{p}$')  
plt.ylabel('Likelihood')  
plt.legend()
```

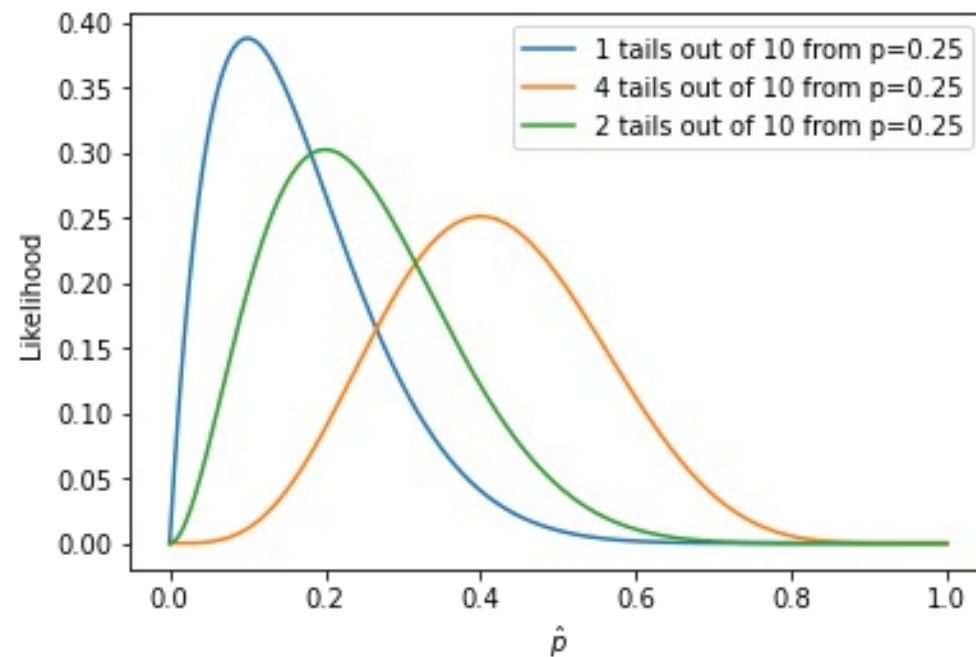
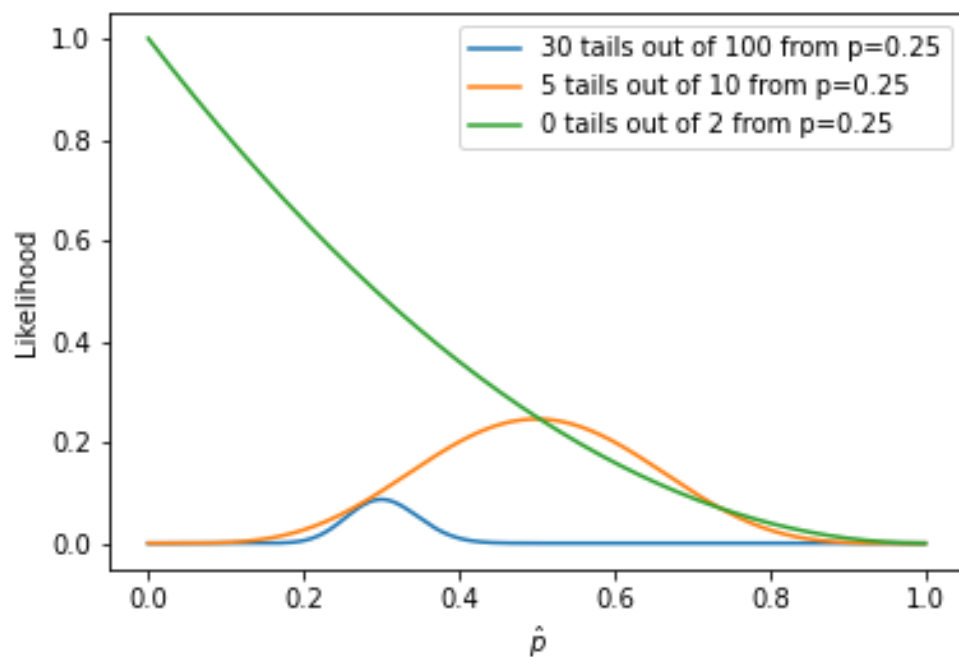


Coin data



`scipy.stats.binom.pmf`

Maximum likelihood estimation



Fair dice data



Roll two dices and observe the numbers on top.

Attributes

(die 1, die 2) $\in \{1, \dots, 6\}^2$

Analysis Questions

What is the distribution of the first die after 100 rolls?

What is the distribution of the sum of the two die?

How likely is the event $A = \{(5,1)\}$?

How likely is the event $B = \{(3,3), (4,3), (4,4), (4,5), (5,4)\}$?

Modified dice data



Roll one die and observe the number on top, reroll the second die until it is equal or greater than the first.

Attributes

(die 1, die 2) $\in \{1, \dots, 6\}^2$

Analysis Questions

What is the distribution of the first die after 100 rolls?


What is the distribution of the sum of the two die?

How likely is the event $A = \{(5,1)\}$?



How likely is the event $B = \{(3,3), (4,3), (4,4), (4,5), (5,4)\}$?

Fair dice data

X is the number rolled on a fair die




X is the sum of numbers rolled on 2 dice



Modified dice data

X is the number rolled on a fair die



X is the sum of numbers rolled on 2 dice

